
A Batch Sequential Halving Algorithm without Performance Degradation

Sotetsu Koyamada
koyamada@atr.jp
ATR, Kyoto University

Soichiro Nishimori
nishimori-soichiro358@g.ecc.u-tokyo.ac.jp
The University of Tokyo

Shin Ishii
ishii@i.kyoto-u.ac.jp
ATR, Kyoto University

Abstract

In this paper, we investigate the problem of pure exploration in the context of multi-armed bandits, with a specific focus on scenarios where arms are pulled in fixed-size batches. Batching has been shown to enhance computational efficiency, but it can potentially lead to a degradation in the performance of the original sequential algorithm due to delayed feedback and reduced adaptability. To address this, we propose a simple batch version of the Sequential Halving algorithm (Karnin et al., 2013) and provide theoretical evidence that batching does not degrade the performance of the original algorithm under practical conditions. Furthermore, we empirically validate our claim through experiments.

1 Introduction

In this study, we consider the pure exploration problem in the field of stochastic multi-armed bandits, which aims to identify the best arm within a given budget (Audibert et al., 2010). Specifically, we concentrate on the *fixed-size batch pulls* setting, where we pull a fixed number of arms simultaneously. Batch computation plays a crucial role in improving computational efficiency, especially in large-scale bandit applications where reward computation can be expensive. For instance, consider applying this to tree search algorithms like Monte Carlo tree search (Tolpin & Shimony, 2012). The reward computation here typically involves the value network evaluation (Silver et al., 2016; 2017), which can be computationally expensive. By leveraging batch computation and hardware accelerators (e.g., GPUs), we can significantly reduce the computational cost of the reward computation. However, while batch computation enhances computational efficiency, its performance (e.g., simple regret) may not match that of sequential computation with the same total budget, due to delayed feedback reducing adaptability. Therefore, the objective of this study is to develop a pure exploration algorithm that maintains its performance regardless of the batch size.

We focus on the *Sequential Halving* (SH) algorithm (Karnin et al., 2013), a popular and well-analyzed pure exploration algorithm. Due to its simplicity, efficiency, and lack of task-dependent hyperparameters, SH finds practical applications in but not limited to hyperparameter tuning (Jamieson & Talwalkar, 2016), recommendation systems (Aziz et al., 2022), and state-of-the-art AlphaZero (Silver et al., 2018) and MuZero (Schrittwieser et al., 2020) family (Danihelka et al., 2022). In this study, we aim to extend SH to a batched version that matches the original SH algorithm’s performance, even with large batch sizes. So far, Jun et al. (2016) introduced a simple batched extension of SH and reported that it performed well in their experiments. However, the theoretical properties of batched SH have not been well-studied in the fixed-size batch pulls setting.

Algorithm 1 SH: Sequential Halving (Karnin et al., 2013)

```
1: input number of arms:  $n$ , total budget:  $T$ 
2: initialize  $\mathcal{S}_0 = [n]$ 
3: for round  $r = 0, \dots, \lceil \log_2 n \rceil - 1$  do
4:   pull each arm  $a \in \mathcal{S}_r$  for  $t_r = \left\lfloor \frac{T}{|\mathcal{S}_r| \lceil \log_2 n \rceil} \right\rfloor$  times
5:    $\mathcal{S}_{r+1} \leftarrow \text{top-}\lceil |\mathcal{S}_r|/2 \rceil$  arms in  $\mathcal{S}_r$  w.r.t. the empirical mean  $\bar{\mu}_a$ 
6: return the only arm in  $\mathcal{S}_{\lceil \log_2 n \rceil}$ 
```

We consider two simple and natural batched variants of SH (Sec. 3): *Breadth-first Sequential Halving* (BSH) and *Advance-first Sequential Halving* (ASH). We introduce BSH as an intermediate step to understanding ASH, which is our main focus. Our main contribution is that we provide a theoretical guarantee for ASH (Sec. 4), showing that *it is algorithmically equivalent to SH as long as the batch budget is not extremely small* — For example, in a 32-armed stochastic bandit problem, ASH can match SH’s choice with 100K sequential pulls using just 20 batch pulls, each of size 5K. This means that ASH can achieve the same performance as SH with significantly fewer pulls when the batch size is reasonably large. Moreover, one can understand the theoretical properties of ASH using the theoretical properties of SH, which have been well-studied (Karnin et al., 2013; Zhao et al., 2023). In our experiments, we validate our claim by comparing the behavior of ASH and SH (Sec. 5.1) and analyze the behavior of ASH with the extremely small batch budget as well (Sec. 5.2).

2 Preliminary

Pure Exploration Problem. Consider a pure exploration problem with n arms. Each arm $a \in [n] := \{1, \dots, n\}$ is associated with an unknown reward distribution characterized by its mean μ_a . Without loss of generality, we assume that $1 \geq \mu_1 \geq \mu_2 \geq \dots \geq \mu_n \geq 0$. In the standard sequential setting, we sequentially pull an arm and observe its stochastic reward for a total budget of T pulls. After these T pulls, we select the arm a_T as the algorithm’s candidate for the best arm. The natural performance measure in pure exploration is the *simple regret*, defined as $\mathbb{E}[\mu_1 - \mu_{a_T}]$ (Bubeck et al., 2009), which compares the performance of the selected arm a_T with the best arm 1.

Sequential Halving (SH; Karnin et al. (2013)) is a sequential elimination algorithm designed for the pure exploration problem. It begins by initializing the set of best arm candidates as $\mathcal{S}_0 = [n]$. In each of the $\lceil \log_2 n \rceil$ rounds, the algorithm halves the set of candidates (i.e., $|\mathcal{S}_{r+1}| = \lceil |\mathcal{S}_r|/2 \rceil$) until it narrows down the candidates to a single arm in $\mathcal{S}_{\lceil \log_2 n \rceil}$. During each round $r \in \{0, \dots, \lceil \log_2 n \rceil - 1\}$, the arms in the active arm set \mathcal{S}_r are pulled equally $t_r := \left\lfloor \frac{T}{|\mathcal{S}_r| \lceil \log_2 n \rceil} \right\rfloor$ times, and the total budget consumed is $T_r := t_r \times |\mathcal{S}_r|$. The SH algorithm is described in Algorithm 1. It has been shown that the simple regret of SH satisfies $\mathbb{E}[\mu_1 - \mu_{a_T}] \leq \tilde{O}(\sqrt{n/T})$, where $\tilde{O}(\cdot)$ ignores the logarithmic factors of n (Zhao et al., 2023). Note that the consumed budget $\sum_{r \leq \lceil \log_2 n \rceil} T_r$ might be less than T . In this study, we assume that the remaining budget is consumed equally by the last two arms.

3 Batch Sequential Halving Algorithms

In this study, we consider the fixed-size batch pulls setting, where we simultaneously pull b arms for B times, with b being the fixed batch size and B being the batch budget (Jun et al., 2016). The standard sequential case corresponds to $b = 1$ and $B = T$. Our objective is to ensure that the performance of the batched variants does not degrade compared to the sequential arm pulls with a budget of $T = b \times B$. In this section, we first reconstruct the SH algorithm so that it can be easily extended to the batched setting (Sec. 3.1). Then, we consider *Breadth-first Sequential Halving* (BSH), one of the simplest batched extensions of SH, as an intermediate step (Sec. 3.2). Finally, we introduce *Advance-first Sequential Halving* (ASH) as a further extension (Sec. 3.3).

Algorithm 2 SH (Karnin et al., 2013) implementation with target pulls $L^{\mathbf{B}}/L^{\mathbf{A}}$

```

1: input number of arms:  $n$ , total budget:  $T$ 
2: initialize empirical mean  $\bar{\mu}_a = 0$  and pull counter  $N_a = 0$  for  $\forall a \in [n]$ 
3: for  $i = 0, \dots, T - 1$  do
4:   let  $\mathcal{A}_i$  be  $\{a \in [n] \mid N_a = L_i\}$   $\triangleright L_i$  is either  $L_i^{\mathbf{B}}$  or  $L_i^{\mathbf{A}}$  (Eq. (1) or (2))
5:    $a_i = \operatorname{argmax}_{a \in \mathcal{A}_i} \bar{\mu}_a$ 
6:   pull arm  $a_i$ : update  $\bar{\mu}_{a_i}$  and  $N_{a_i} \leftarrow N_{a_i} + 1$ 
7: return  $\operatorname{argmax}_{a \in [n]} (N_a, \bar{\mu}_a)$   $\triangleright$  best arm in active arms

```

3.1 SH implementation with target pulls

As BSH/ASH is a natural batched extension of SH, we first reconstruct the implementation of the SH algorithm as Algorithm 2 so that it can be easily extended to BSH/ASH. Note that, in this study, the operation $\operatorname{argmax}_{x \in \mathcal{X}} (\ell_x, m_x)$ selects the element $x \in \mathcal{X}$ that firstly maximizes ℓ_x . If multiple elements achieve this maximum, it then selects among these the one that maximizes m_x . At the i -th arm pull, SH selects the arm a_i that has the highest empirical reward $\bar{\mu}_a$ among the candidates \mathcal{A}_i :

$$a_i = \operatorname{argmax}_{a \in \mathcal{A}_i} \bar{\mu}_a,$$

where $\mathcal{A}_i := \{a \in [n] \mid N_a = L_i\}$ is the candidates at i -th arm pull, N_a is the total number of pulls of arm a , and L_i is the *target pull* defined as either **breadth-first** manner

$$L_i^{\mathbf{B}} := \underbrace{\sum_{r' < r(i)} t_{r'}}_{\text{pulls before } r(i)} + \underbrace{\left\lfloor \frac{i - \sum_{r' < r(i)} T_{r'}}{|\mathcal{S}_{r(i)}|} \right\rfloor}_{\text{pulls in } r(i)}, \quad (1)$$

or **advance-first** manner

$$L_i^{\mathbf{A}} := \underbrace{\sum_{r' < r(i)} t_{r'}}_{\text{pulls before } r(i)} + \underbrace{\left(\left(i - \sum_{r' < r(i)} T_{r'} \right) \bmod t_{r(i)} \right)}_{\text{pulls in } r(i)}, \quad (2)$$

where $r(i)$ is the round of the i -th arm pull. This $L_i^{\mathbf{B}}/L_i^{\mathbf{A}}$ represents the cumulative number of pulls of the arm selected at the i -th pull before the i -th arm pull. We omitted the dependency on n and T for simplicity. The definition of $L_i^{\mathbf{B}}/L_i^{\mathbf{A}}$ is a bit complicated, and it may be straightforward to write down the algorithm that constructs $L^{\mathbf{B}} := \{L_0^{\mathbf{B}}, \dots, L_T^{\mathbf{B}}\}$ and $L^{\mathbf{A}} := \{L_0^{\mathbf{A}}, \dots, L_T^{\mathbf{A}}\}$ as shown in Algorithm. 3 and Algorithm. 4, respectively. Note that the choice of $L_i^{\mathbf{B}}/L_i^{\mathbf{A}}$ is arbitrary and does not affect the behavior of SH — as long as the arm pull is sequential and not batched. Python code for this SH implementation is available in App. A. Note that using target pulls to implement SH is natural and not new. For example, Mctx¹ (Babuschkin et al., 2020) has a similar implementation.

3.2 BSH: Breadth-first Sequential Halving

Now, we extend SH to BSH, in which we select arms so that the number of pulls of each arm becomes as equal as possible using $L^{\mathbf{B}}$. When pulling arms in a batch, we need to consider not only the number of pulls of the arms but also the number of scheduled pulls in the current batch. Therefore, we introduce *virtual pull count* M_a , the number of scheduled pulls of arm a in the current batch. For i -th batch pull, we sequentially select b arms with the highest empirical rewards from the candidates $\{a \in [n] \mid N_a + M_a = L_i^{\mathbf{B}}\}$ and pull them in the batch. BSH algorithm is described in App. B. BSH is similar to a batched extension of SH introduced in Jun et al. (2016) in the sense that it selects arms so that the number of pulls of each arm becomes as equal as possible.

¹<https://github.com/google-deepmind/mctx>

Algorithm 3 Breadth-first target pulls L^{B}

```
1: input number of arms:  $n$ , total budget:  $T$ 
2: initialize empty  $L^{\text{B}}$ ,  $K = n$ ,  $T_r = 0$ 
3: for  $r = 0, \dots, \lceil \log_2 n \rceil - 1$  do
4:   for  $j = 0, \dots, t_r - 1$  do
5:     for  $k = 0, \dots, K - 1$  do
6:       append  $T_r + j$  to  $L^{\text{B}}$ 
7:    $K \leftarrow \lceil K/2 \rceil$  and  $T_r \leftarrow T_r + t_r$ 
8: return  $L^{\text{B}}$ 
```

Algorithm 4 Advance-first target pulls L^{A}

```
1: input number of arms:  $n$ , total budget:  $T$ 
2: initialize empty  $L^{\text{A}}$ ,  $K = n$ ,  $T_r = 0$ 
3: for  $r = 0, \dots, \lceil \log_2 n \rceil - 1$  do
4:   for  $k = 0, \dots, K - 1$  do
5:     for  $j = 0, \dots, t_r - 1$  do
6:       append  $T_r + j$  to  $L^{\text{A}}$ 
7:    $K \leftarrow \lceil K/2 \rceil$  and  $T_r \leftarrow T_r + t_r$ 
8: return  $L^{\text{A}}$ 
```

Algorithm 5 ASH: Advance-first Sequential Halving

```
1: input number of arms:  $n$ , batch size:  $b$ , batch budget:  $B$ 
2: initialize empirical mean  $\bar{\mu}_a = 0$  and pull counter  $N_a = 0$  for  $\forall a \in [n]$ 
3: for  $i = 0, \dots, B - 1$  do
4:   initialize empty batch  $\mathcal{B}_j$  and virtual pulls  $M_a = 0$  for  $\forall a \in [n]$ 
5:   for  $j = 0, \dots, b - 1$  do
6:     let  $\mathcal{A}_i$  be  $\{a \in [n] \mid N_a + M_a = L_{i \times b + j}^{\text{A}}\}$   $\triangleright L_{i \times b + j}^{\text{A}}$  uses  $T = b \times B$ 
7:      $a_i = \operatorname{argmax}_{a \in \mathcal{A}_i} (N_a, \bar{\mu}_a)$ 
8:     push  $a_i$  to  $\mathcal{B}_j$  and  $M_{a_i} \leftarrow M_{a_i} + 1$ 
9:   batch pull arms in  $\mathcal{B}_j$   $\triangleright$  batch execution may be efficient
10:  update  $\bar{\mu}_a$  and  $N_a \leftarrow N_a + M_a$  for all  $a \in \mathcal{B}_j$ ,
11: return  $\operatorname{argmax}_{a \in [n]} (N_a, \bar{\mu}_a)$ 
```

3.3 ASH: Advance-first Sequential Halving

We further extend SH to ASH in a similar manner as BSH. Algorithm 5 shows the ASH algorithm. The difference between BSH and ASH is that

1. ASH selects arms in *advance-first* manner using L^{A} instead of L^{B} (line 6), and
2. ASH considers not only the empirical rewards $\bar{\mu}_a$ but also the number of actual pulls N_a when selecting arms in a batch (line 7).

The second difference is necessary to ensure that the arm with the highest empirical reward is selected from the arms that have actually finished pulling in the batch spanning two rounds. In Sec. 4, we will show that ASH is algorithmically equivalent to SH with the same total budget $T = b \times B$. Python code for ASH is available in App. A.

4 Algorithmic Equivalence of SH and ASH

This section presents a theoretical guarantee for the ASH algorithm.

Theorem 1 *Given a stochastic bandit problem with $n \geq 2$ arms, let $b \geq 2$ be the batch size and B be the batch budget satisfying $B \geq \max\{4, \frac{n}{b}\} \lceil \log_2 n \rceil$. Then, the ASH algorithm (Algorithm 5) is algorithmically equivalent to the SH algorithm (Algorithm 2) with the same total budget $T = b \times B$ — when the two algorithms are executed under the same pseudo-random number, the selected arms are always the same.*

Proof. The condition $B \geq \max\{4, \frac{n}{b}\} \lceil \log_2 n \rceil$ is divided into two separate conditions:

$$B \geq \frac{n}{b} \lceil \log_2 n \rceil, \tag{C1}$$




Figure 1: Visualization of SH (Karnin et al., 2013), BSH (Sec. 3.2), and ASH (Sec. 3.3) for 8-armed bandit problem. For SH, total budget $T = 24 \times 8 = 192$. For BSH and ASH, batch size $b = 24$ and batch budget $B = 8$. The same color indicates the same batch pull.

and

$$B \geq 4 \lceil \log_2 n \rceil. \quad (\text{C2})$$

A key observation is that the ASH and SH algorithms differ only when a batch pull spans two rounds. Thus, we focus on the scenario where a batch pull spans two rounds. In this case, let $z < b$ be the number of pulls that consume the budget for round r , and $b - z$ be the number of pulls that consume the budget for round $r + 1$. The following proposition is demonstrated: $\forall n \geq 2, \forall b \geq 2, \forall r < \lceil \log_2 n \rceil - 1, \forall z < b$, if (C1) and (C2) hold, then

$$|\mathcal{S}_{r+1}| - \left\lceil \frac{b-z}{t_{r+1}} \right\rceil \geq \left\lceil \frac{z}{t_r} \right\rceil. \quad (3)$$

The left-hand side (LHS) of Eq. (3) denotes the number of arms progressing to the subsequent round post-batch pull, whereas the right-hand side (RHS) quantifies the arms pending completion of their pulls at the batch pull juncture. This inequality, if satisfied, ensures that arms supposed to advance to the next round in SH are not left behind in ASH even when a batch spans two rounds. Demonstrating the scenario where $z = b - 1$ suffices, as it represents the worst-case condition. Let $x := |\mathcal{S}_r| \geq 3$ for the given $r < \log_2 n - 1$. Two cases are considered. **Case 1:** when $n \leq 4b$. Given that $t_r = \lfloor \frac{b \times B}{x \lceil \log_2 n \rceil} \rfloor \geq \lfloor 4b/x \rfloor$ as derived from (C2), it is sufficient to show

$$\left\lceil \frac{x}{2} \right\rceil - 1 \geq \left\lceil \frac{b-1}{\lfloor 4b/x \rfloor} \right\rceil \quad (4)$$

in $x \in [3, 4b]$. This assertion is directly supported by Lemma 1. **Case 2:** when $4b < n$. Given that $t_r = \lfloor \frac{b \times B}{x \lceil \log_2 n \rceil} \rfloor \geq \lfloor n/x \rfloor$ as derived from (C1), it is sufficient to show $\left\lceil \frac{x}{2} \right\rceil - 1 \geq \left\lceil \frac{n/4-1}{\lfloor n/x \rfloor} \right\rceil$ in $x \in [3, n]$. This conclusion follows by the same reasoning applied in Case 1. This completes the proof.

Lemma 1 For any integer $b \geq 2$, the inequality $\left\lceil \frac{x}{2} \right\rceil - 1 \geq \left\lceil \frac{b-1}{\lfloor 4b/x \rfloor} \right\rceil$ holds for all integers $x \in [3, 4b]$.

The proof of Lemma 1 is in App. C. Here, we provide the visualization of Eq. (4) in Fig. 2 to intuitively show that Lemma 1 holds. Each colored line represents the RHS for different $b \leq 32$. One can see that the LHS is always greater than the RHS for any $x \in [3, 4b]$.

Remark 1 The condition (C1) is common to both SH and ASH — SH implicitly assumes $T \geq n \lceil \log n \rceil$ as the minimum condition to execute. This is because we need to pull each arm at least once in the first round (i.e., $t_1 \geq 1$). With the same argument, the batch budget B must satisfy (C1). On the other hand, (C2) is specific to ASH and is required to ensure the equivalence. As we discuss in the Sec. 4.1, we claim that this additional (C2) is not practically problematic.




Figure 2: Lemma 1.

Remark 2 Note that the condition (C2) is tight; Theorem 1 does not hold even if $B \geq \alpha \lceil \log_2 n \rceil$ for any positive value $\alpha < 4$.

Proof. We aim to demonstrate the existence of a value x such that $\lceil \frac{x}{2} \rceil - 1 - \frac{b-1}{\lceil \alpha b/x \rceil} < 0$ when $n \leq \alpha b$. Consider the case when $x = 4$. In this scenario, the LHS of the inequality can be expressed as $1 - \frac{b-1}{\lceil \alpha b/4 \rceil} < 1 - \frac{4}{\alpha} \frac{b-1}{b} \rightarrow 1 - \frac{4}{\alpha}$ as $b \rightarrow \infty$. As $\alpha < 4$, it follows that LHS < 0 for sufficiently large values of b . This completes the proof.

Remark 3 Theorem 1 implies that, unlike BSH, ASH inherits the theoretical properties of SH, including the simple regret bound (Zhao et al., 2023) and that ASH can arbitrarily increase the batch size b (and decrease the batch budget B) without performance degradation to improve the computational efficiency as long as the condition (C1) and (C2) hold.

4.1 Discussion on the conditions

To show that SH and ASH are algorithmically equivalent, we used an additional condition (C2) of $\mathcal{O}(\log n)$. However, we claim that this condition is not practically problematic because the condition (C1), the minimum condition required to execute (unbatched) SH, is dominant ($\mathcal{O}(n \log n)$). This condition (C1) is dominant over (C2) as shown in Fig. 3. We can see that the condition (C2) only affects the algorithm when the batch size is sufficiently larger than the number of arms ($b \gg n$). This is a reasonable result, meaning that we cannot guarantee the equivalent behavior to SH with an extremely small batch budget, such as $B = 1$. On the other hand, if the user secures the minimum budget $B = 4 \lceil \log_2 n \rceil$ that depends only on the number of arms n and increases only logarithmically, regardless of the batch size b , they can increase the batch size arbitrarily and achieve the same result as when SH is executed sequentially with the same total budget, with high computational efficiency.




Figure 3: ■ both (C2) and (C1) hold (i.e., ASH is equivalent to SH). ■ only (C2) holds (i.e., SH is not executable). ■ only (C1) holds (i.e., SH is executable but ASH may not be equivalent to SH). ■ neither (C2) nor (C1) holds.

5 Empirical Validation

We conducted experiments to empirically demonstrate that ASH maintains its performance for large batch size b , in comparison to its sequential counterpart SH. To evaluate this, we utilized a polynomial family parameterized by α as a representative batch problem instance, where the reward gap $\Delta_i := \mu_1 - \mu_i$ follows a polynomial distribution with parameter α : $\Delta_i \propto (i/n)^\alpha$ (Jamieson et al., 2013; Zhao et al., 2023). This choice is motivated by the observation that real-world applications exhibit polynomially distributed reward gaps, as mentioned in Zhao et al. (2023). In our study, we considered three different values of α (0.5, 1.0, and 2.0) to capture various reward distributions (see Fig. 4). Additionally, we characterized each bandit problem instance by specifying the minimum and maximum rewards, denoted as μ_{\min} and μ_{\max} respectively. Hence, we denote a bandit problem instance as $\mathcal{T}(n, \alpha, \mu_{\min}, \mu_{\max})$.

We also implemented a simple batched extension of SH introduced by Jun et al. (2016) as a baseline for comparison. We refer to this algorithm as Jun+16. The implementation of Jun+16 is described




Figure 4: Polynomial(α)




Figure 5: Single regret comparison of BSH, ASH, and Jun+16 against SH when $B \geq 4\lceil \log_2 n \rceil$.




Figure 6: Single regret comparison of BSH, ASH, and Jun+16 against SH when $B < 4\lceil \log_2 n \rceil$.

in App. D. Jun et al. (2016) did not provide a theoretical guarantee for Jun+16, but it has shown comparable to or better performance than their proposed algorithm in their experiments.

5.1 Large batch budget scenario: $B \geq 4\lceil \log_2 n \rceil$

First, we empirically confirm that, as we claimed in Sec. 4, ASH is indeed equivalent to SH under the condition (C2). We generated 10K instances of bandit problems and applied ASH and SH to each instance with different 100 seeds. We randomly sampled n from $\{2, \dots, 1024\}$, α from $\{0.5, 1.0, 2.0\}$, and μ_{\min} and μ_{\max} from $\{0.1, 0.2, \dots, 0.9\}$. For each instance $\mathcal{T}(n, \alpha, \mu_{\min}, \mu_{\max})$, we randomly sampled the batch budget $B \leq 10\lceil \log_2 n \rceil$ and the batch size $b \leq 5n$ so that the condition (C1) and (C2) are satisfied. As a result, we confirmed that the selected arms of ASH and SH are identical in all 10K instances and 100 seeds for each instance. We also conducted the same experiment for BSH and Jun+16. We plotted the simple regret of BSH, ASH, and Jun+16 against SH in Fig. 5. There are 10K instances, and each point represents the average simple regret of 100 seeds for each instance. To compare the performance, we fitted a linear regression model to the simple regret of BSH, ASH, and Jun+16 against SH as $y = \beta x$, where y is the simple regret of BSH, ASH, or Jun+16, x is the simple regret of SH. The slope β is estimated by the least squares method. The estimated slope β is 1.008 for BSH, 1.000 for ASH, and 0.971 for Jun+16, which indicates that the simple regret of ASH, BSH, and Jun+16 is almost the same as SH on average.

5.2 Small batch budget scenario: $B < 4\lceil \log_2 n \rceil$

Next, we examined the performances of BSH, ASH, and Jun+16 against SH when the condition (C2) is not satisfied, i.e., when the batch budget is extremely small $B < 4\lceil \log_2 n \rceil$ and thus Theorem 1 does not hold. We conducted the same experiment as in Sec. 5.1 except the batch budget $B < 4\lceil \log_2 n \rceil$. We sampled B so that B is larger than the number of rounds. The results are shown in Fig. 6. The slope β is estimated as 1.059 for BSH, 1.011 for ASH, and 1.017 for Jun+16. All the estimated slopes are worse than when $B \geq 4\lceil \log_2 n \rceil$. However, the estimated slopes are still close to 1, which indicates that while we do not have a theoretical guarantee, the performance of BSH, ASH, and Jun+16 is comparable to SH on average.

6 Related Work

Sequential Halving. Among the algorithms for the pure exploration problem in multi-armed bandits (Audibert et al., 2010), Sequential Halving (SH; Karnin et al. (2013)) is one of the most

popular algorithms. The theoretical properties of SH have been well studied (Karnin et al., 2013; Zhao et al., 2023). Due to its simplicity, SH has been widely used for these (but not limited to) applications: In the context of *tree-search* algorithms, As the root node selection of Monte Carlo tree search can be regarded as a pure exploration problem (Tolpin & Shimony, 2012), Danihelka et al. (2022) incorporated SH into the root node selection and significantly reduced the number of simulations to improve the performance during AlphaZero/MuZero training. From the min-max search perspective, some studies recursively applied SH to the internal nodes of the search tree (Cazenave, 2014; Pepels et al., 2014). SH is also used for *hyperparameter optimization*; Jamieson & Talwalkar (2016) formalized the hyperparameter optimization problem in machine learning as a *non-stochastic* multi-armed bandit problem, where the reward signal is not from stochastic stationary distributions but from deterministic function changing over training steps. Li et al. (2018; 2020) applied SH to hyperparameter optimization in asynchronous parallel settings, which is similar to our batch setting. Their asynchronous approach may have *incorrect promotions* to the next rounds but is more efficient than the synchronous approach. Aziz et al. (2022) applied SH to *recommendation systems*, in which identifies appealing podcasts for users.

Batched bandit algorithms. Batched bandit algorithms have been studied in various contexts (Perchet et al., 2016; Gao et al., 2019; Esfandiari et al., 2021; Jin et al., 2021a;b; Kalkanli & Ozgur, 2021; Karbasi et al., 2021; Provodin et al., 2022). Among the batched bandit studies for the pure exploration problem (Agarwal et al., 2017; Grover et al., 2018; Jun et al., 2016), Jun et al. (2016) is the most relevant to our work as they also consider the *fixed-size batch pulls* setting. To the best of our knowledge, the first batched SH variant with a fixed batch size b was introduced by Jun et al. (2016) as a baseline algorithm in their study (Jun+16). It is similar to BSH and it pulls arms so that the number of pulls of the arms is as equal as possible (breadth-first manner). They reported that Jun+16 experimentally performs comparable to or better than their proposed method but did not provide a theoretical guarantee for Jun+16. Our ASH is different from their batch variant in that ASH pulls arms in an advance-first manner instead of a breadth-first manner.

7 Limitation and Future Work

Our batched variants of SH assume that the reward distributions of the arms are from i.i.d. distributions. This property is essential to allow batch pulls. One limitation is that it may be difficult to apply our algorithms to bandit problems where the reward distribution is non-stationary. For example, Jamieson & Talwalkar (2016) applied SH to hyperparameter tuning, where rewards are time-series losses during model training. We cannot apply our batched variants to this problem because we cannot observe “future losses” in a batch.

Our batched variants of SH are suitable for tasks where we can evaluate arms efficiently by pulling them in a batch rather than sequentially. For example, if the evaluation of arms depends on the output of neural networks and can be evaluated quickly in a batch by utilizing accelerators such as GPUs (e.g., Monte Carlo tree search with value network as Danihelka et al. (2022)). Applying our batched variants to such algorithms is a possible future direction.

8 Conclusion

In this paper, we proposed ASH as a simple and natural extension of the SH algorithm. We theoretically showed that ASH is algorithmically equivalent to SH as long as the batch budget is not excessively small. This allows ASH to inherit the well-studied theoretical properties of SH, including the simple regret bound. Our experimental results confirmed this claim and demonstrated that ASH and other batched variants of SH like Jun+16 perform comparably to SH in terms of simple regret. These findings suggest that we can utilize simple batched variants of SH for efficient evaluation of arms with large batch sizes while avoiding performance degradation compared to the sequential execution of SH. By providing a practical solution for efficient arm evaluation, our study opens up

new possibilities for applications that require large budgets. Overall, our work highlights the batch robust nature of SH and its potential for large-scale bandit problems.

References

- Arpit Agarwal, Shivani Agarwal, Sepehr Assadi, and Sanjeev Khanna. [Learning with Limited Rounds of Adaptivity: Coin Tossing, Multi-Armed Bandits, and Ranking from Pairwise Comparisons](#). In *COLT*, 2017.
- Jean-Yves Audibert, Sébastien Bubeck, and Rémi Munos. [Best Arm Identification in Multi-armed Bandits](#). In *COLT*, 2010.
- Maryam Aziz, Jesse Anderton, Kevin Jamieson, Alice Wang, Hugues Bouchard, and Javed Aslam. [Identifying new podcasts with high general appeal using a pure exploration infinitely-armed bandit strategy](#). In *RecSys*, 2022.
- Igor Babuschkin, Kate Baumli, Alison Bell, Surya Bhupatiraju, Jake Bruce, Peter Buchlovsky, David Budden, Trevor Cai, Aidan Clark, Ivo Danihelka, Claudio Fantacci, Jonathan Godwin, Chris Jones, Ross Hemsley, Tom Hennigan, Matteo Hessel, Shaobo Hou, Steven Kapturowski, Thomas Keck, Iurii Kemaev, Michael King, Markus Kunesch, Lena Martens, Hamza Merzic, Vladimir Mikulik, Tamara Norman, John Quan, George Papamakarios, Roman Ring, Francisco Ruiz, Alvaro Sanchez, Rosalia Schneider, Eren Sezener, Stephen Spencer, Srivatsan Srinivasan, Luyu Wang, Wojciech Stokowiec, and Fabio Viola. [The DeepMind JAX Ecosystem](#). <https://github.com/google-deepmind>, 2020.
- Sébastien Bubeck, Rémi Munos, and Gilles Stoltz. [Pure Exploration in Multi-armed Bandits Problems](#). In *ALT*, 2009.
- Tristan Cazenave. [Sequential Halving Applied to Trees](#). *IEEE T-CIAIG*, 7(1):102–105, 2014.
- Ivo Danihelka, Arthur Guez, Julian Schrittwieser, and David Silver. [Policy improvement by planning with Gumbel](#). In *ICLR*, 2022.
- Hossein Esfandiari, Amin Karbasi, Abbas Mehrabian, and Vahab Mirrokni. [Regret Bounds for Batched Bandits](#). In *AAAI*, 2021.
- Zijun Gao, Yanjun Han, Zhimei Ren, and Zhengqing Zhou. [Batched Multi-armed Bandits Problem](#). In *NeurIPS*, 2019.
- Aditya Grover, Todor Markov, Peter Attia, Norman Jin, Nicolas Perkins, Bryan Cheong, Michael Chen, Zi Yang, Stephen Harris, William Chueh, and Stefano Ermon. [Best arm identification in multi-armed bandits with delayed feedback](#). In *AISTATS*, 2018.
- Kevin Jamieson and Ameeet Talwalkar. [Non-stochastic Best Arm Identification and Hyperparameter Optimization](#). In *AISTATS*, 2016.
- Kevin Jamieson, Matthew Malloy, Robert Nowak, and Sebastien Bubeck. [On Finding the Largest Mean Among Many](#). *arXiv:1306.3917*, 2013.
- Tianyuan Jin, Jing Tang, Pan Xu, Keke Huang, Xiaokui Xiao, and Quanquan Gu. [Almost Optimal Anytime Algorithm for Batched Multi-Armed Bandits](#). In *ICML*, 2021a.
- Tianyuan Jin, Pan Xu, Xiaokui Xiao, and Quanquan Gu. [Double Explore-then-Commit: Asymptotic Optimality and Beyond](#). In *COLT*, 2021b.
- Kwang-Sung Jun, Kevin Jamieson, Robert Nowak, and Xiaojin Zhu. [Top Arm Identification in Multi-Armed Bandits with Batch Arm Pulls](#). In *AISTATS*, 2016.
- Cem Kalkanli and Ayfer Ozgur. [Batched Thompson Sampling](#). In *NeurIPS*, 2021.

-
- Amin Karbasi, Vahab Mirrokni, and Mohammad Shadravan. [Parallelizing Thompson Sampling](#). In *NeurIPS*, 2021.
- Zohar Karnin, Tomer Koren, and Oren Somekh. [Almost Optimal Exploration in Multi-Armed Bandits](#). In *ICML*, 2013.
- Liam Li, Kevin Jamieson, Afshin Rostamizadeh, Ekaterina Gonina, Jonathan Ben-tzur, Moritz Hardt, Benjamin Recht, and Ameet Talwalkar. [A System for Massively Parallel Hyperparameter Tuning](#). In *MLSys*, 2020.
- Lisha Li, Kevin Jamieson, Giulia DeSalvo, Afshin Rostamizadeh, and Ameet Talwalkar. [Hyperband: A Novel Bandit-Based Approach to Hyperparameter Optimization](#). *JMLR*, 18(185):1–52, 2018.
- Tom Pepels, Tristan Cazenave, Mark HM Winands, and Marc Lanctot. [Minimizing Simple and Cumulative Regret in Monte-Carlo Tree Search](#). In *CGW*, 2014.
- Vianney Perchet, Philippe Rigollet, Sylvain Chassang, and Erik Snowberg. [Batched bandit problems](#). *Ann. Stat.*, 44(2):660 – 681, 2016.
- D. Provodin, P. Gajane, M. Pechenizkiy, and M. Kaptein. [The Impact of Batch Learning in Stochastic Linear Bandits](#). In *ICDM*, 2022.
- Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, et al. [Mastering Atari, Go, chess and shogi by planning with a learned model](#). *Nature*, 588(7839):604–609, 2020.
- David Silver, Aja Huang, Chris J Maddison, Arthur Guez, Laurent Sifre, George van den Driessche, Julian Schrittwieser, Ioannis Antonoglou, Veda Panneershelvam, Marc Lanctot, Sander Dieleman, Dominik Grewe, John Nham, Nal Kalchbrenner, Ilya Sutskever, Timothy Lillicrap, Madeleine Leach, Koray Kavukcuoglu, Thore Graepel, and Demis Hassabis. [Mastering the game of Go with deep neural networks and tree search](#). *Nature*, 529(7587):484–489, 2016.
- David Silver, Julian Schrittwieser, Karen Simonyan, Ioannis Antonoglou, Aja Huang, Arthur Guez, Thomas Hubert, Lucas Baker, Matthew Lai, Adrian Bolton, Yutian Chen, Timothy Lillicrap, Fan Hui, Laurent Sifre, George van den Driessche, Thore Graepel, and Demis Hassabis. [Mastering the game of Go without human knowledge](#). *Nature*, 550(7676):354–359, 2017.
- David Silver, Thomas Hubert, Julian Schrittwieser, Ioannis Antonoglou, Matthew Lai, Arthur Guez, Marc Lanctot, Laurent Sifre, Dhharshan Kumaran, Thore Graepel, Timothy Lillicrap, Karen Simonyan, and Demis Hassabis. [A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play](#). *Science*, 362(6419):1140–1144, 2018.
- David Tolpin and Solomon Shimony. [MCTS Based on Simple Regret](#). In *AAAI*, 2012.
- Yao Zhao, Connor Stephens, Csaba Szepesvari, and Kwang-Sung Jun. [Revisiting Simple Regret: Fast Rates for Returning a Good Arm](#). In *ICML*, 2023.

A Python code

For the sake of reproducibility and better understanding, we provide Python code for the Sequential Halving (SH) algorithm using target pulls and Advance-first Sequential Halving (ASH) algorithm in Fig. 7.

```
from math import log2, ceil, floor
import numpy as np

def sh(bandit: BanditProblem, n: int, T: int) -> int:
    L = _get_target_pulls(n, T)          # L: target pulls
    N = np.append(np.zeros(n, dtype=int), -1e9) # N: pull counts
    R = np.append(np.zeros(n, dtype=float), 0.) # R: avg rewards
    for i in range(T):
        a = np.argmax(np.where(N == L[i], R, -np.inf))
        r = bandit.pull(a)
        R[a] = (R[a] * N[a] + r) / (N[a] + 1)
        N[a] += 1
    return int(np.argmax(np.where(N >= max(N), R, -np.inf)))

def ash(bandit: BanditProblem, n: int, B: int, b: int = 1) -> int:
    L = _get_target_pulls(n, b * B)      # L: target pulls
    N = np.append(np.zeros(n, dtype=int), -1e9) # N: pull counts
    R = np.append(np.zeros(n, dtype=float), 0.) # R: avg rewards
    for i in range(B):
        batch = []
        M = np.zeros_like(N)              # M: virtual pull counts
        for j in range(b):
            N_max = np.max(np.where(N + M == L[i * b + j], N, -np.inf))
            a = np.argmax(np.where((N + M == L[i * b + j]) & (N == N_max), R, -np.inf))
            batch.append(a)
            M[a] += 1
        rewards = bandit.batch_pull(batch)
        for a, r in zip(batch, rewards):
            R[a] = (R[a] * N[a] + r) / (N[a] + 1)
            N[a] += 1
    return int(np.argmax(np.where(N >= max(N), R, -np.inf)))

def _get_target_pulls(n: int, T: int) -> list[int]:
    target_pulls = []
    num_rounds = ceil(log2(n))
    num_active_arms = n
    cum_pulls = 0
    for r in range(num_rounds):
        t_r = floor(T / (num_active_arms * num_rounds))
        if r == num_rounds - 1:
            remaining_pulls = T - len(target_pulls)
            t_r = remaining_pulls // 2
        for _ in range(num_active_arms):
            for i in range(t_r):
                target_pulls.append(cum_pulls + i)
            cum_pulls += t_r
        num_active_arms = ceil(num_active_arms / 2) # halving
    return target_pulls + [int(-1e9)] * (T - len(target_pulls))
```

Figure 7: Python implementation of the SH algorithm using target pulls (Algorithm. 2) and ASH algorithm (Algorithm. 5).

B BSH algorithm

Algorithm. 6 shows the detailed BSH algorithm Sec. 3.2.

Algorithm 6 BSH: Breadth-first Sequential Halving

```

1: input number of arms:  $n$ , batch size:  $b$ , batch budget:  $B$ 
2: initialize empirical mean  $\bar{\mu}_a$  and pull counter  $N_a = 0$  for  $\forall a \in [n]$ 
3: for  $i = 0, \dots, B - 1$  do
4:   initialize empty batch  $\mathcal{B}_j$  and virtual pulls  $M_a = 0$  for  $\forall a \in [n]$ 
5:   for  $j = 0, \dots, b - 1$  do
6:     let  $\mathcal{A}_i$  be  $\{a \in [n] \mid N_a + M_a = L_{i \times b + j}^{\mathbf{B}}\}$   $\triangleright L_{i \times b + j}^{\mathbf{B}}$  uses  $T = b \times B$ 
7:      $a_i = \operatorname{argmax}_{a \in \mathcal{A}_i} \bar{\mu}_a$ 
8:     push  $a_i$  to  $\mathcal{B}_j$  and  $M_{a_i} \leftarrow M_{a_i} + 1$ 
9:   batch pull arms in  $\mathcal{B}_j$   $\triangleright$  batch execution may be efficient
10:  update  $\bar{\mu}_a$  and  $N_a \leftarrow N_a + M_a$  for all  $a \in \mathcal{B}_j$ ,
11: return  $\operatorname{argmax}_{a \in [n]} (N_a, \bar{\mu}_a)$ 

```

C Proof of Lemma 1

Lemma 1 For any integer $b \geq 2$, the inequality

$$\left\lceil \frac{x}{2} \right\rceil - 1 \geq \left\lceil \frac{b-1}{4b/x} \right\rceil \quad (5)$$

holds for all integers $x \in [3, 4b]$.

Proof. This proof demonstrates that for any integer $b \geq 2$ and $x \in [3, 4b]$, the inequality (5) is satisfied. Given $z \geq c \implies z \geq \lceil c \rceil$ for any integer z and real number c , it suffices to demonstrate that

$$\left\lceil \frac{x}{2} \right\rceil - 1 \geq \frac{b-1}{4b/x} \iff \left\lceil \frac{x}{2} \right\rceil - 1 - \frac{b-1}{4b/x} \geq 0.$$

Given that $\lfloor \frac{4b}{x} \rfloor > 0$, it follows that

$$\left(\left\lceil \frac{x}{2} \right\rceil - 1 \right) \left\lfloor \frac{4b}{x} \right\rfloor - (b-1) \geq 0, \quad (6)$$

for any integer $b \geq 2$ and $x \in [3, 4b]$. Two cases are considered:

Case 1: x is even. Suppose $x = 2y$, with $y \in [2, 2b]$. It is aimed to demonstrate that

$$(y-1) \left\lfloor \frac{2b}{y} \right\rfloor - (b-1) \geq 0. \quad (7)$$

Two sub-cases are considered:

1. For $y \in [b+1, 2b]$, as $\left\lfloor \frac{2b}{y} \right\rfloor = 1$, LHS $= (y-1) - (b-1) \geq 0$.
2. For $y \in [2, b]$, as $\lfloor c \rfloor \geq c-1$ for any real number c , we have LHS $\geq (y-1) \left(\frac{2b}{y} - 1 \right) - (b-1) = -\frac{(y-2)(y-b)}{y}$. As $y > 0$ and $-(y-2)(y-b) \geq 0$ in $y \in [2, b]$, we have LHS ≥ 0 .

Consequently, it has been established that for even values of x , the inequality (7) is upheld.

Case 2: x is odd. Suppose $x = 2y + 1$, with $y \in [1, 2b - 1]$. It is aimed to demonstrate that

$$y \left\lfloor \frac{4b}{2y+1} \right\rfloor - (b-1) \geq 0. \quad (8)$$

Two sub-cases are considered:

1. For $y \in [b, 2b - 1]$, as $\left\lfloor \frac{4b}{2y+1} \right\rfloor = 1$, LHS = $y - (b - 1) \geq 0$.
2. For $y \in [1, b - 1]$, as $\lfloor c \rfloor \geq c - 1$ for any real number c , we have LHS $\geq y \left(\frac{4b}{2y+1} - 1 \right) - (b - 1) = \frac{2by - b - 2y^2 + y + 1}{2y+1} = \frac{-2y(y - (b + \frac{1}{2})) - (b - 1)}{2y+1} \geq 0$. As $2y + 1 > 0$ and $-2y(y - (b + \frac{1}{2})) - (b - 1) \geq 0$ in $y \in [1, b - 1]$, we have LHS ≥ 0 .

Similarly, it has been demonstrated that for odd values of x , the inequality (8) is upheld.

Therefore, through the analysis of these two cases, it is proven that for any integer $b \geq 2$ and $x \in [3, 4b]$, the inequality (6) is satisfied, thereby confirming the validity of (5). This completes the proof.

D Batch Sequential Halving introduced in Jun et al. (2016)

[Algorithm. 7](#) shows the detailed batch version of the Sequential Halving algorithm introduced in Jun et al. (2016).

Algorithm 7 Batched Sequential Halving introduced in Jun et al. (2016)

- 1: **input** number of arms: n , batch budget: B , batch size: b
 - 2: **initialize** $\mathcal{S}_0 \leftarrow \{1, \dots, n\}$
 - 3: **for** round $r = 0, \dots, \lceil \log_2 n \rceil - 1$ **do**
 - 4: **for** $1, \dots, \left\lfloor \frac{B}{\lceil \log_2 n \rceil} \right\rfloor$ **do**
 - 5: select batch actions \mathcal{B} so that the number of pulls of each arm in \mathcal{S}_r is as equal as possible
 - 6: pull arms \mathcal{B} in the batch
 - 7: $\mathcal{S}_{r+1} \leftarrow \text{top-}\lceil |\mathcal{S}_r|/2 \rceil$ arms in \mathcal{S}_r w.r.t. empirical rewards
 - 8: **return** the only arm in $\mathcal{S}_{\lceil \log_2 n \rceil}$
-